



NODUS · COMPLETE VAULT GUIDE

Databases

Build approachable relational datasets, inspect quality and analyse patterns.

The Databases vault provides typed tables, relations, formulas, rollups, views, search, analysis and grounded data chat without requiring a separate database server.

EDITION	English
VERSION	Nodus 4.1
UPDATED	14 August 2026
AUDIENCE	Researchers, collections teams and anyone whose project has outgrown a spreadsheet.

Contents

Core operation, configuration and the complete vault workflow.

Core Nodus workflow	3
1 Welcome to Nodus	3
2 Create and switch vaults	4
3 Navigation and global search	5
4 Library, sources and reader	5
5 AI providers and task models	6
6 Privacy, backups and recovery	7
7 Toolkit and integrations	7
Complete Databases workflow	9
1 Home and database catalogue	9
2 Columns and field types	10
3 Rows and record detail	11
4 Relations, formulas and rollups	11
5 Table, gallery and saved views	12
6 Filter, sort and bulk actions	13
7 CSV import and export	14
8 Cross-database search	14
9 Analysis and charts	15
10 Data chat	15
11 AI-assisted columns	16
12 Database-linked notes	17
13 Database settings and safety	18
Completion checklist	19

PART I · FOUNDATION

Core Nodus workflow

These operating principles apply to every vault. They establish the local-first data boundary, source workflow, optional model configuration and recovery discipline used throughout this manual.

START HERE

1 Welcome to Nodus

Nodus is a local-first desktop workspace for research, teaching and study. Work is organised into focused vaults that share one private engine.


A vault changes the navigation, workflows and assistant context without changing the core promises: your database and indexes remain on your computer, accounts are optional, and AI runs only when a feature needs it and you have configured a provider.

STEP BY STEP

1. Install the current package for macOS, Windows or Linux from the Nodus release page.
2. Open Nodus and choose the interface language, Nodi style and tutorial path.
3. Create a vault for the kind of work you are doing, or load its sample vault before importing your own material.

GOOD PRACTICE

- Use one vault per coherent project, course, family or dataset.
- Sample vaults are disposable and are the safest way to explore.



English interface shown with the built-in sample vault.

START HERE

2 Create and switch vaults

Create independent workspaces and move between them without mixing their context.

The vault switcher is the boundary between projects. Each vault has its own notes, analysis and domain records, while the global Library can link the same source into compatible vaults without duplicating it.

STEP BY STEP

1. Open the vault switcher from the application header and choose New vault.
2. Name the vault, choose its type and review the short description before confirming.
3. Use the same switcher to move to another vault; Nodus restores the last open section.

GOOD PRACTICE

- Use explicit names such as 'MA thesis - memory' or 'Biology 2026'.
- Changing a vault type after substantial work is not a substitute for creating the correct vault.

ESSENTIALS

3 Navigation and global search

The sidebar exposes the workflow for the active vault; search retrieves material across the relevant local index.

Sections are grouped by intent - explore, analyse and write, or organisation and assessment. Search combines exact keyword matching with semantic retrieval when an embedding model is configured, and results preserve links to their source.

STEP BY STEP

1. Choose a section in the left sidebar; the active section is highlighted.
2. Open Search and enter at least two meaningful characters or a natural-language concept.
3. Filter or inspect a result, then follow its source link to the document, record, person or note.

GOOD PRACTICE

- Search for concepts, not only exact titles.
- Rebuild an index from Settings if recently imported content does not appear.

ESSENTIALS

4 Library, sources and reader

Keep preserved originals beside clean reading copies, citations, highlights and extracted evidence.


The cross-vault Library can mirror Zotero, import RIS, BibTeX and CSL JSON, or accept files directly. Nodus preserves the original attachment and creates searchable derivatives. Compatible sources can be linked to several vaults without copying the original.

STEP BY STEP

1. Open Library and add a local file, import bibliographic data or connect Zotero.
2. Place the item in a collection and add useful tags; wait for extraction and indexing to finish.
3. Open the item, choose the clean reading copy or preserved original, then highlight passages and link the item into the active vault.

GOOD PRACTICE

- Keep the preserved original for page-accurate verification.
- Treat automatically extracted text and metadata as a first pass that may need correction.



English interface shown with the built-in sample vault.

CONFIGURATION

5 AI providers and task models

Bring your own cloud key or use compatible local models, with a different model for each task.

Provider keys are stored in the operating-system keychain. Nodus distinguishes deterministic local tools from model-assisted features and shows when content will leave the device. Model preferences can mix providers by extraction, chat, writing and report tasks.

STEP BY STEP

1. Open Settings, then Providers, and configure a cloud key or local endpoint.
2. Use the validation action before saving, then open Models.
3. Assign a suitable model to each task and run a small test on sample content.

GOOD PRACTICE

- A local model improves privacy but may require more memory and time.
- Review every generated claim against its linked evidence.

CONFIGURATION

6 Privacy, backups and recovery

Protect local work with a verified backup and understand exactly when network services are used.


Vault databases, indexes and library files are local by default. Backups can include all vaults and the shared Library. Nodus reports backup health and creates protected recovery copies before major migrations.

STEP BY STEP

1. Open Settings, choose a backup destination and enable automatic backups.
2. Run a manual backup, then use verification to confirm that its manifest and files are readable.
3. Review the privacy notice for each optional online feature before sending content.

GOOD PRACTICE

- Store backups on a different physical device or a trusted encrypted location.
- Do not treat sync as a backup; keep versioned recovery copies.



English interface shown with the built-in sample vault.

CONFIGURATION

7 Toolkit and integrations

Convert, protect, translate, OCR and present documents, or connect Zotero, Word, MCP clients and Nodus Server.

Toolkit utilities are available across vaults. Integrations remain optional: Zotero supplies bibliographic material, Word uses grounded Nodus context, MCP exposes approved local tools, and Nodus Server publishes a selected copy rather than the source vault.

STEP BY STEP

1. Open Toolkit and choose Convert, Protect, Translate, OCR Workspace or PDF Presenter.
2. Select a file from disk or a compatible vault source and review the output options.
3. For an integration, enable only the connector you need and test it with non-sensitive sample content first.


GOOD PRACTICE

- Permanent redaction in Protect is local and irreversible in the exported copy.
- Published server spaces should contain only material intended for their readers.

PART II · VAULT GUIDE

Complete Databases workflow

The Databases vault provides typed tables, relations, formulas, rollups, views, search, analysis and grounded data chat without requiring a separate database server.



English interface shown with the built-in sample vault.

ORIENTATION

1 Home and database catalogue

Create, open and monitor the databases in a vault.


Home lists databases, row totals and shortcuts to search, analysis and data chat.

STEP BY STEP

1. Create an empty database or import an existing CSV.
2. Give the database a specific name and description.
3. Open it from the catalogue or sidebar.

GOOD PRACTICE

- Use separate databases for distinct entity types, then connect them with relations.



English interface shown with the built-in sample vault.

BUILD

2 Columns and field types

Model text, numbers, dates, options, files, relations, formulas and AI-assisted outputs explicitly.


Typed fields make filtering, validation and analysis more reliable than a sheet of ambiguous strings. Schema changes should reflect the meaning of the data.

STEP BY STEP

1. Add a column and choose the narrowest useful field type.
2. Name it in plain language and configure options, precision or validation.
3. Test the field on several records before bulk import.

GOOD PRACTICE

- Do not encode several facts in one text column if they must be filtered separately.



English interface shown with the built-in sample vault.

BUILD

3 Rows and record detail

Edit records in the table or open a full record card with attachments and long fields.

Each row is a record. The detail view makes dense schemas easier to inspect and supports file attachments and extended content.

STEP BY STEP

1. Add a row and fill its identifying field first.
2. Edit compact values inline; open the record for long text, files or a complete audit.
3. Review required or invalid values before leaving the record.

GOOD PRACTICE

- Choose one stable field as the human-readable record title.

BUILD

4 Relations, formulas and rollups

Connect records and calculate derived values without copying data.


Relations link a row to records in another database. Rollups summarise related values, while formulas compute deterministic outputs from the current record.

STEP BY STEP

1. Create a relation column and choose its target database and cardinality.
2. Link existing records rather than retyping their names.
3. Add a rollup or formula, then test null values and edge cases.

GOOD PRACTICE

- Prefer relations over duplicated text when the referenced entity can change.



English interface shown with the built-in sample vault.

EXPLORE

5 Table, gallery and saved views

Present the same records in layouts and filters suited to different tasks.

Views do not duplicate data. They store layout, visible fields, filters and sort order so a catalogue, review queue and fieldwork view can coexist.

STEP BY STEP

1. Create or switch to Table or Gallery.

2. Choose visible fields and configure filter groups and sort order.
3. Save the view with a task-oriented name.

GOOD PRACTICE

- Keep one unfiltered audit view so hidden records are never mistaken for missing data.

EXPLORE

6 Filter, sort and bulk actions

Narrow records with explicit conditions and apply reviewed changes efficiently.

Filters can combine conditions; sorts establish a stable order. Bulk operations are powerful and should follow a preview or backup.

STEP BY STEP

1. Add a filter condition, operator and value; group conditions deliberately.
2. Add one or more sort keys and check how blanks are ordered.
3. Select records, preview the bulk action and confirm only the intended fields.

GOOD PRACTICE

- Export or back up before irreversible bulk changes.

The screenshot displays the Nodus Databases demo interface. The top navigation bar includes links for Vaults, How it works, Views, Wiki, Contribute, and FAQ, along with a language selector (English), a star icon, a download icon (633 downloads), and a 'Try the live demo' button. The main content area shows a table titled 'Field samples' with 8 entries (100%). The table columns are: Name, Code, Species, Habitat, Weight, Cover, and Collection Date. The data rows are as follows:

Name	Code	Species	Habitat	Weight	Cover	Collection Date
Alpine moss	MC-001	Moss	Mountain Forest	12.4	38	2026-01-01
Rock lichen	MC-002	Lichen	Mountain	8.1	22	2026-01-01
Shade fern	MC-003	Fern	Forest River	20	55	2026-01-01
Green alga	MC-004	Alga	Coast	5.7	12	2026-01-01
Riverbank moss	MC-005	Moss	River Forest	9.9	31	2026-01-01
Coastal lichen	MC-006	Lichen	Coast	3.2	9	2026-01-01
Royal fern	MC-007	Fern	Forest	27.6	61	2026-01-01
Brown alga	MC-008	Alga	Coast	14.3	27	2026-01-01

The interface also includes a sidebar with navigation options: Home, EXPLORE (Search, Field samples, Experiments, Reading list), ANALYZE (Analysis, Data chat), and WRITE (Notes, Settings). A 'Web demo - Databases mode' notice is visible at the top of the main content area.

English interface shown with the built-in sample vault.

EXCHANGE

7 CSV import and export

Bring in existing tabular data and export CSV, XLSX or JSON.

Import previews headers and type inference before writing rows. Export respects the selected database and format but may differ in how it represents relations or files.

STEP BY STEP

1. Clean obvious duplicate headers, then open Import and preview encoding and delimiter.
2. Map each source column to a field and review inferred dates, numbers and options.
3. After import, check row count and samples; before export, choose the scope and target format.

GOOD PRACTICE

- Keep the original import file unchanged as provenance.

EXPLORE

8 Cross-database search

Find database names and matching content inside rows.


Search spans every database in the vault and returns the record context needed to open the correct row.

STEP BY STEP

1. Enter a distinctive name, identifier or phrase.
2. Use database grouping to disambiguate similar records.
3. Open the row and verify the value in context.

GOOD PRACTICE

- Normalise identifiers and names to improve reliable retrieval.



English interface shown with the built-in sample vault.

ANALYSE

9 Analysis and charts

Inspect distributions, missing values, summaries and visual patterns.

Deterministic statistics and charts run from the selected database. Model-assisted reports are optional and should cite the fields and populations used.

STEP BY STEP

1. Choose a database and review row count, field completeness and detected types.
2. Select a categorical or numeric field for an appropriate chart.
3. Investigate outliers and missingness in the underlying rows before interpreting a pattern.

GOOD PRACTICE

- A chart is only as trustworthy as the schema and data-entry rules behind it.

ANALYSE

10 Data chat

Ask grounded questions about a selected database and receive inspectable summaries.


Data chat scopes answers to database rows rather than general knowledge. Responses may include computed summaries or charts, but important results should be checked against deterministic analysis.

STEP BY STEP

1. Select the target database.
2. Ask a precise question that names the measure, grouping and population.
3. Inspect cited or referenced records and reproduce critical totals in Analysis or export.

GOOD PRACTICE

- Ask for assumptions and missing-data treatment explicitly.



English interface shown with the built-in sample vault.

ANALYSE

11 AI-assisted columns

Classify or summarise records into a new field under human review.

An AI column applies a defined instruction to selected record fields. Output is a proposal and can vary with model or prompt changes.

STEP BY STEP

1. Create an AI-assisted field and define its input columns and expected output format.

2. Run it on a small sample and refine the instruction.
3. Review results before processing the full set and record the model configuration.

GOOD PRACTICE

- Use deterministic formulas for calculations; reserve AI fields for interpretive tasks.

WRITE

12 Database-linked notes

Document definitions, cleaning decisions and interpretations beside the records.


Notes can link to databases and rows, providing an audit trail for decisions that do not fit in the schema itself.

STEP BY STEP

1. Create a note for a method, issue or finding.
2. Link the affected database or records.
3. Record the decision, date and consequence for later analysis.

GOOD PRACTICE

- Maintain a data dictionary and change log from the beginning.



English interface shown with the built-in sample vault.

CONFIGURE

13 Database settings and safety

Control import behaviour, models, integrations and recoverable storage.

Settings combine local database maintenance with shared model, MCP and backup controls.

STEP BY STEP

1. Review provider and model choices for analysis and AI columns.
2. Test MCP access with a read-only client before exposing write operations.
3. Enable automatic backups and export critical datasets periodically.

GOOD PRACTICE

- Treat schema migrations and bulk edits as high-risk changes deserving a backup.

REFERENCE

Completion checklist

- The vault scope and name match the project.
- Sources or records have preserved originals and reviewed metadata.
- Generated interpretations have been checked against their evidence.
- Optional providers and integrations use only the intended data.
- Exports have been proofread in their final format.
- A recent backup has been created and verified.

This manual documents the stable English workflow represented in Nodus 4.1. Interface details may evolve between releases; the in-app privacy notice and release notes govern the exact behaviour of the installed version.